



UNIVERSITY OF ALBERTA

DEPARTMENT OF
COMPUTING SCIENCE

An Introduction to Berkeley DB

CMPUT 291

**File and Database Management
Systems**

Objectives

- Levels of abstraction in a typical DBMS.
- Introduction to Berkeley DB.
 - What is Berkeley DB.
 - File Organizations supported by Berkeley DB.
 - Access method operations, cursors.
- Take a look at the Berkeley DB API using simple examples in Python.



Levels of abstraction in a typical DBMS

- Data is stored in tables in conceptual schema.
- High level SQL commands (DDL or DML statements) are used to manipulate data in database.
- Data is actually organized in files and indexes in physical schema.
- A DBMS uses suitable storage structure, index files, and access paths to evaluate queries and return results.



What Is Berkeley DB?


- It is an open source embedded database library that provides a simple function-call API for data access and management.
- It supports a number of file structures and operations.
- You can use it to:
 - (1) Develop programs working with storage structures and indexes.
 - (2) Develop applications that don't need the full functionality of a DBMS and require high performance.



File organizations supported in Berkeley DB

- **Hash:**
 - ✓ Data is stored in an extended linear hash table
 - ✓ Good for applications that need quick random-access.
- **Btree:**
 - ✓ Data is stored in a sorted, balanced tree structure
 - ✓ Good for range-based searches

For more information see [BerkeleyDB Access Methods](#)



File organizations supported in Berkeley DB (Cont.)

Record-number-based:

The logical record number is the primary key of records (compared to Hash and Btree organizations when one or a combination of attributes form a key for records)

- **Queues:**

- ✓ Data is stored in a queue as fixed-length records
- ✓ Good for fast inserts at the tail of the queue
- ✓ It also supports read/delete operation from the head of the queue.
- ✓ provides record-level locking (as opposed to the page-level locking that the other access methods use)

- **Recnos:**

- ✓ Data is stored in either fixed or variable-length records.


For more information see [BerkeleyDB Access Methods](#)



A summary of methods in Database class

Method	Description
open	Create/Open a database
close	Close a database
get	Get items from a database
put	Store items into a database
delete	Delete items from a database
stat	Return database statistics
truncate	Emptying a database of all records
sync	Flushes all modified records from the DB cache to disk
associate	Declares one DB as secondary index of a primary DB

For more information see [BerkeleyDB API](#)



A summary of methods in Cursor class

- A database cursor supports traversing the database.
- It is the only way to access individual duplicate records.

Method	Description
dup	Duplicates a cursor
close	Discards the cursor
count	Returns the number of duplicate data items for the key to which the cursor refers
delete	Deletes the key/value pair to which the cursor refers
get	Retrieves key/value pairs from the database
set	Moves the cursor to the specified key in the database and return the associated key/value pair
put	Stores key/value pairs into the database using a cursor

For more information see [BerkeleyDB API](#)



Note: Connecting to the lab machine using ssh

<https://eclass.srv.ualberta.ca/mod/page/view.php?id=3105765>

Berkeley DB API (Python)

1. Create the database:

```
from bsddb3 import db

database = db.DB() #handle for Berkeley DB database
DB_File = "fruit.db"
database.open(DB_File ,None, db.DB_HASH, db.DB_CREATE)
# database.open(DB_File ,None, db.DB_BTREE, db.DB_CREATE)
# database.open(DB_File ,None, db.DB_QUEUE, db.DB_CREATE)
# database.open(DB_File ,None, db.DB_RECNO, db.DB_CREATE)
# The arguments correspond to (fileName, database name within the file for multiple
databases, database type, flag to create database)
```



Berkeley DB API (Python)

2. Declare the cursor for the database:

```
curs = database.cursor()
```

3. Insert records into the database:

```
# insertion of (key, value) pair using the cursor:  
#The arguments corresponds to (key, value, flag to insert key-value pair )  
curs.put(b"apple", "red ", db.DB_KEYFIRST)  
  
# insertion using the database object's put method  
database.put(b"pear", "green")
```



Berkeley DB API (Python)

4. Iterate through the database to display all key-value pairs:

```
# use the cursor
iter = curs.first()
while iter:
    print(iter)
    iter = curs.next()
```



Berkeley DB API (Python)

5. Retrieve a specific key-value or index-value pair

```
# using the database object's get method: only retrieves the value
```

```
result = database.get(b'pear')
```

```
print(result)
```

```
# b'green'
```

```
# using the cursor object's set method:
```

```
result = curs.set(b'pear')
```

```
#Moves the cursor to the specified key in the database and return the associated key/value pair.
```

```
print(result)
```




Berkeley DB API (Python)

6. Remove a (key,value)

```
# remove using the cursor – deletes the key-value pair currently referenced by the cursor:  
curs.current() #Returns the key/value pair currently referenced by the cursor  
curs.delete()  
# our database now has (apple,red) only  
# remove by using database object  
database.delete(b'apple')  
# Now our database is empty.
```

7. Close database and cursor:

```
curs.close()  
database.close()
```



Berkeley DB API (Python)

8. Show records with the same key (duplicates):

```
# Set duplicate flag before you create the database:
# db.DB_DUP allows duplicate data items in the tree
database.set_flags(db.DB_DUP)

#Duplicate key-value pairs could be traversed using cursor

# prints no. of k-v pairs that have the same key (for the key which the cursor is pointing to)
print(curs.count())

# prints the next k-v pair if it is a duplicate
print(curs.next_dup())

#returns 'None' if the next k-v pair of the database is not a duplicate data record for the
current key/value pair
```



Exercise 1: Iterate_All

- ❖ Write a code to iterate through all key-value pairs(including the duplicates) using next_dup() method for duplicate keys.
 - In order to test your code, add the following key-value pairs to see if they are retrieved:

```
database.put(b'key1', "value1")  
database.put(b'key1', "value2")  
database.put(b'key2', "value1")  
database.put(b'key2', "value2")
```



Exercise 1: Iterate_All (Cont.)

- ❖ Iterating through k-v pair including duplicates:

```
curs = database.cursor()
iter = curs.first()

while (iter):

    print(curs.count())
    #prints no. of rows that have the same key for the current key-value pair referred by cursor.

    print(iter)

    #iterating through duplicates:
    dup = curs.next_dup()
    while(dup!=None):
        print(dup)
        dup = curs.next_dup()

    iter = curs.next()
```

- ❖ A solution is provided on e-class (named Iterate_All.py)



Exercise 2: Populate bdb

Write a program that:

- Gets a student name(as a key) from input
 - Searches the database if there is any key(student) with the same name
 - Prints any found key as well as its value
 - Then, asks the user if the input key should be inserted, and if yes, the program asks for its value as well
(let's say the value is the student's mark)
 - Terminates if "q" is entered.
- Pay attention that duplicate keys should be supported (multiple marks for a student)
- Hint:
- ✓ Convert Char String to Byte String: `name.encode("utf-8")`
 - ✓ Convert Byte String to Char String: `value.decode("utf-8")`



Exercise 2: Populate bdb (Cont.)

- ❖ In order to set the cursor on the first data item for the input name:

```
name = input("Enter a student Name to look up: ")  
result = curs.set(name.encode("utf-8"))
```

- ❖ To get rid of the byte prefix (b' ') for printing students' names and marks:

```
print("Name: "+str(result[0].decode("utf-8"))+", Mark: "+str(result[1].decode("utf-8")))
```

- ❖ A solution is provided on e-class (named Populate_bdb.py)



Exercise 3: Range Search


Write a program that:

- Gets a name called `Starting_Name`
- Gets another name called `Ending_Name`
- Searches for all students' names from the previous exercise that come after or in the same position as `Starting_Name` (when sorted alphabetically) and come before `Ending_Name`.
- Prints the names and marks of the students found in the previous step.



Hint:

- ✓ In Btree DB type, By default, the sort order is lexicographical, with shorter keys coming before longer keys.
- ✓ `Curs.set_range(key)`: In Btree DB type, sets the cursor to key/value pair which is the smallest key greater than or equal to the specified key, and returns the pair.



Exercise 3: Range Search

- ❖ To get the record that has the smallest key greater than or equal to the Starting Name:

```
Starting_Name = input("Enter the Starting_Name: ")  
result = curs.set_range(Starting_Name.encode("utf-8"))
```

- ❖ To check if the student's name comes before Ending_Name:

```
if (str(result[0].decode("utf-8"))[0:len(Ending_Name)]) < Ending_Name):  
    ...
```

- ❖ A solution is provided on e-class (named Range_Search.py)
- ❖ **Question:** What if the keys were **numbers** instead of names, and you were asked to output the k-v pairs with keys **higher** than **Starting_Number** and **Lower** than **Ending_Number**?



What's Next?

- Check out the page “[introduction to Python bsddb3](#)” for more information.
- Take a look at Berkeley DB API for C & Java in this [tuto rial](#).

