

Embedded SQL

Davood Rafiei



Introduction

- **Basic Idea:** Use SQL statements inside a host language (C, C++, Java, Python, ...).
- **Advantages:**
 - Can do all the fancy things you do in C/Java/Python.
 - Still have the power of SQL.
- **Need mechanisms for**
 - Embedding SQL statements in code
 - Transferring data to/from DBMS
 - Compiling and linking the program.



SQL inside Applications

- Statement-level interface
 - SQL statements appear as new statements in the program.
 - Precompile step: replaces new statements with some procedure calls.
 - Flavors: **Static SQL**, **Dynamic SQL**
- Call-level interface
 - Program is written entirely in the host language
 - No precompile step.
 - SQL statements are passed as parameters to some procedures.
 - Flavors: **ODBC**, **JDBC**, **sqlite callback**



Statement-level Interface

- New statements:

- EXEC SQL <sql statement>
- E.g.
 - ✓ EXEC SQL SELECT ...
 - ✓ EXEC SQL UPDATE ...

- Static SQL: e.g.

```
EXEC SQL SELECT COUNT(*) INTO :cnt
        FROM emp;
```

The SQL statement is known at compile time.

- Dynamic SQL: e.g.

```
strcpy(qstr, "SELECT COUNT(*) FROM emp");
EXEC SQL PREPARE Q FROM :qstr;
```

The SQL statement is not known at compile time.



Program Structure

SQL Include statements

```
main(int argc, char *argv[])
```

```
{
```

Declarations



Connect to the Database



Do your work with the database



Process Errors



Commit/Rollback

```
}
```



SQL Include Statements

- An essential one: SQL Communication Area
 - **#include <sqlca.h>** , or
 - **EXEC SQL INCLUDE SQLCA**
- It provides runtime status information including:
 - sqlca.sqlcode: the return code of SQL statements
 - sqlca.sqlerrm: the error messages (if any)
 - ✓ sqlca.sqlerrm.sqlerrmc : error message
 - ✓ sqlca.sqlerrm.sqlerrml : the length of the error message



Host Variables

- Used to pass values between a SQL statement and the rest of the program.
- Declaration: as follows

```
EXEC SQL BEGIN DECLARE SECTION;  
/* all host variable declarations go here */  
int    cnt;  
char name[20];  
EXEC SQL END DECLARE SECTION;
```
- Usage in SQL statements:

```
EXEC SQL SELECT COUNT(*) INTO :cnt FROM emp;
```

 - The variable must be preceded with `:' to distinguish it from SQL identifiers.



Error Handling

- Check `sqlca.sqlcode`, the return code of SQL statements
 - `== 0` : the command was successful
 - `> 0` : no row in the output
 - `< 0` : error
- Use **WHenever** statements: e.g.

```
EXEC SQL WHENEVER SQLERROR      GOTO lbl1;  
EXEC SQL WHENEVER NOT FOUND    GOTO lbl3;
```

```
EXEC SQL WHENEVER SQLERROR      DO err_func();
```

```
EXEC SQL WHENEVER SQLERROR      CONTINUE;
```



Error Handling

- Example:

```
EXEC SQL DROP TABLE emp;
if (sqlca.sqlcode == 0) {
    printf("Table emp is successfully dropped\n");
}
else {
    sqlca.sqlerrm.sqlerrmc[sqlca.sqlerrm.sqlerrml] = '\0';
    printf("Oracle error: %s\n", sqlca.sqlerrm.sqlerrmc);
}
```



myprog.pc: Our First Program

```
#include <stdio.h>
EXEC SQL INCLUDE SQLCA;

int main(int argc, char *argv[])
{
    EXEC SQL BEGIN DECLARE SECTION;
        char user[10] = "SCOTT";
        char pwd[10] = "TIGER";
    EXEC SQL END DECLARE SECTION;

    EXEC SQL WHENEVER SQLERROR GOTO error;
    EXEC SQL CONNECT :user IDENTIFIED BY :pwd;
    printf("connected to Oracle");

    /* SQL statements */

    EXEC SQL COMMIT RELEASE;
    return 0;
```



myprog.pc (Cont.)

error:

```
sqlca.sqlerrm.sqlerrmc[sqlca.sqlerrm.sqlerrml] = '\0';  
printf("Oracle Error: %s\n", sqlca.sqlerrm.sqlerrmc);
```

```
EXEC SQL WHENEVER SQLERROR CONTINUE;  
EXEC SQL ROLLBACK RELEASE;
```

```
return 1;
```

```
}
```



Program Preparation

- Precompile
proc myprog.pc
 - result: myprog.c (a pure C program with SQL statements replaced with library calls)
- Compile the C program
cc myprog.c
 - result: myprog.o
- Link the libraries
cc -o myprog myprog.o -L...
 - result: myprog (an executable program)
- ** Use a makefile instead **
 - Check the Oracle directory for a sample *makefile*



myprog2.pc: Our 2nd Program

```
/* get the emp_number from input and print the emp_name */
```

```
...
```

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
    int emp_number;
```

```
    char emp_name[30];
```

```
EXEC SQL END DECLARE SECTION;
```

```
EXEC SQL WHENEVER SQLERROR    GOTO error;
```

```
EXEC SQL WHENEVER NOT FOUND    GOTO nope;
```

```
/* Connect to the database ..... */
```

```
printf("Enter emp_number:");
```

```
scanf("%d", emp_number);
```

```
EXEC SQL SELECT ename INTO :emp_name
```

```
        FROM emp
```

```
        WHERE empno = :emp_number;
```

```
printf("Employee name is %s\n", emp_name);
```

```
return 0;
```

```
error: ...
```

```
nope: ...
```



Complications

- What if a host variable is assigned a NULL?
 - not a constant.
 - solution: use an extra **indicator variable**.
- What if a SQL statement returns more than one row?
 - cannot fit it in any C variable!
 - solution: use a **cursor**.



Indicator Variables

```
/* get the emp_number from input and print the emp_name */  
...
```

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
    int emp_number;
```

```
    char emp_name[30];
```

```
    short emp_name_ind;
```

```
EXEC SQL END DECLARE SECTION;
```

```
...
```

```
EXEC SQL SELECT ename INTO :emp_name INDICATOR :emp_name_ind  
        FROM   emp  
        WHERE  empno = :emp_number;
```

```
if (emp_name_ind < 0)
```

```
    printf("Employee name is NULL\n");
```

```
else
```

```
    printf("Employee name is %s\n", emp_name);
```

```
...
```



Use of a Cursor

```
/* print the names of all employees */  
...  
EXEC SQL BEGIN DECLARE SECTION;  
    char emp_name[30];  
EXEC SQL END DECLARE SECTION;  
  
EXEC SQL DECLARE emp_cursor CURSOR FOR  
        SELECT ename  
        FROM   emp;  
  
EXEC SQL OPEN emp_cursor;  
EXEC SQL WHENEVER NOT FOUND GOTO end;  
  
printf("Employee names are:\n");  
for (;;) {  
    EXEC SQL FETCH emp_cursor INTO :emp_name;  
    printf("%s\n", emp_name);  
}  
end:  
EXEC SQL CLOSE emp_cursor;  
EXEC SQL COMMIT RELEASE;  
return 0;
```




Cursor Overview

- ***Result set*** – set of rows produced by a SELECT statement
- ***Cursor*** – pointer to a row in the result set.
- Cursor operations:
 - *Declaration*
 - *Open* – execute SELECT to determine result set and initialize pointer
 - *Fetch* – advance pointer and retrieve next row
 - *Close* – deallocate cursor



Cursor




Statement-level Interface (Summary)

- One database at a time
- Both schema (for input and output) and database name must be known at compile time
- Ordered from the most efficient to the least
 - Static SQL
 - Dynamic SQL
 - Call-level interface: JDBC



Call-level Interface

- Neither the schema nor the database is known at compile time.
- The application can connect to more than one database.



Executing a Query

```
import java.sql.*;    // imports all classes in package java.sql
String db_url = "jdbc:sqlite:/Users/drafiei/291dbfile.db";
String driverName = "org.sqlite.JDBC";
//String db_url = "jdbc:oracle:thin:@gwynne.cs.ualberta.ca:1521:CRS";
//String driverName = "oracle.jdbc.driver.OracleDriver";
```

```
Class.forName(driverName);    // loads the specified driver
```

```
Connection con = DriverManager.getConnection(db_url)
//Connection con = DriverManager.getConnection(db_url, Id, Pwd);
```

- connects to the DBMS at address *db_url*
- If successful, creates a connection object, *con*, for managing the connection

```
Statement stat = con.createStatement ();
```

- Creates a statement object *stat*
- Statements have *executeQuery()* method



Executing a Query

```
String query = "SELECT  T.StudId FROM Transcript T" +  
               "WHERE  T.CrsCode = 'cse305' " +  
               "AND  T.Semester = 'S2000' ";
```

```
ResultSet res = stat.executeQuery(query);
```

- *Creates a result set object, res.*
- *Prepares and executes the query.*
- *Stores the result set produced by execution in res (analogous to opening a cursor).*
- *The query string can be constructed at run time (as above).*
- *The input parameters are plugged into the query when the string is formed (as above)*



Result Sets and Cursors

- Three types of result sets in JDBC:
 - *Forward-only*: not scrollable
 - *Scroll-insensitive*: scrollable (can jump up and down in the result set!), changes made to underlying tables after the creation of the result set are not visible through that result set
 - *Scroll-sensitive*: scrollable, changes made to the tuples in a result set after the creation of that set are visible through the result set



Result Set

```
Statement stat = con.createStatement (
    ResultSet.TYPE_SCROLL_SENSITIVE,
    ResultSet.CONCUR_UPDATABLE );
```

- Any result set type can be declared *read-only* or *updatable* – **CONCUR_UPDATABLE** (assuming SQL query satisfies the conditions for updatable views)
- Current row of an updatable result set can be updated or deleted, and a new row can be inserted, causing changes in base table

```
res.updateString ("Name", "John" );    // update attribute "Name" of
                                         // current row in row buffer.
```

```
or < res.updateRow ( );                // make the change permanent
     res.CancelRowUpdate ( );          // cancel the update
```



SQLite callback in C

- Three functions: open, exec, close
 - `sqlite3_open(const char *filename, sqlite3 **db)`
 - `sqlite3_exec(sqlite3 *db, const char *sql, sqlite_callback, void *data, char **errmsg)`
 - `sqlite3_close(sqlite3 *db)`
- Output is processed by a callback function



myp.c

```
##include <stdio.h>
#include <stdlib.h>
#include <sqlite3.h>

static int callback(void *data, int argc, char **argv, char **aColName){
    int i;
    fprintf(stderr, "%s: ", (const char*)data);

    for(i = 0; i<argc; i++){
        printf("%s = %s\n", aColName[i], argv[i] ? argv[i] : "NULL");
    }

    printf("\n");
    return 0;
}
```



```

int main(int argc, char* argv[]) {
    sqlite3 *db;  char *zErrMsg = 0;
    int rc;  char *sql;
    const char* data = "Callback function called";

    /* Open database */
    rc = sqlite3_open("291lect.db", &db);
    if( rc ) {
        fprintf(stderr, "Can't open database: %s\n", sqlite3_errmsg(db));
        return(0);
    } else {
        fprintf(stderr, "Opened database successfully\n");
    }

    /* Create and execute SQL statement */
    sql = "SELECT * from customer";
    rc = sqlite3_exec(db, sql, callback, (void*)data, &zErrMsg);

    if( rc != SQLITE_OK ) {
        fprintf(stderr, "SQL error: %s\n", zErrMsg);
        sqlite3_free(zErrMsg);
    } else {
        fprintf(stdout, "Operation done successfully\n");
    }
    sqlite3_close(db);
    return 0;
}

```



Compile and Run

- Compile

```
gcc -Wall -std=c99 -L/usr/lib/sqlite3 myp.c -lsqlite3 -o myp
```

- Run

```
./myp
```

Opened database successfully

Callback function called: cname = Davood

street = 114 St

city = Edmonton

Callback function called: cname = Ehsan

street = University Ave

city = NULL

Operation done successfully

More information
https://www.sqlite.org/c_interface.html



More Information

- Statement-level Interface:
 - Not supported in SQLite
 - But is supported in other databases such as Oracle
- Call-level Interface
 - JDBC tutorials
 - Callback interface in C



Summary

- Covered:
 - Static SQL
 - Call level interface (JDBC, C)
- SQLite in Python is covered in the lab
- Not Covered: Dynamic SQL, ODBC
- Final note (but quite important):
 - Avoid data processing in the host language if it can be passed to SQL.
 - Use the host language for things that cannot be done in SQL.

