

# SQL - Basics

Davood Rafiei

Original material Copyright 2001-2018  
(some material from textbooks and other instructors)

# SQL Overview

- **Structured Query Language**
  - standard query language for relational system
  - developed in IBM Almaden (system R)
- **Some features**
  - Declarative: specify the properties that should hold in the result, not how to obtain the result
  - Complex queries have procedural elements
  - Set/Bag semantics
- **International Standards**
  - SQL1 (1986)
  - SQL2 (SQL-92)
  - SQL3 (SQL-99)
  - Later extensions in 2003, 2006, 2008, 2011, 2016

# SQL Overview (cont)

- Data definition component
  - CREATE TABLE table-name (col-defs, constraints)
  - DROP TABLE table-name
  - ALTER TABLE table-name action
    - ✓ modifies the definition of a table where action is
      - ✓ ADD (col-def)
      - ✓ MODIFY (col-def)
      - ✓ ADD constraint
      - ✓ etc.
- Data update component
  - INSERT INTO table-name ...
  - DELETE FROM table-name ...
  - UPDATE table-name ...
- Query Language

# Example Tables

**branch** (bname, address, city, assets)

**customer**(cname, street, city)

**deposit**(accno, cname, bname, balance)

**loan**(accno, cname, bname, amount)

```
CREATE TABLE branch (  
    bname      CHAR(20),  
    address    CHAR(30),  
    city       CHAR(15),  
    asset      NUMBER(38)  
);
```

Oracle command

(in Oracle, a semicolon  
always marks the end of a  
SQL statement.)

# Simple Queries

**branch** (bname, address, city, assets)

- Find the names of all branches with assets greater than \$2,500,000.

```
SELECT bname  
FROM    branch  
WHERE    assets > 2500000
```

- Find the names of all branches in Edmonton with assets greater than \$2,500,000.

```
SELECT bname  
FROM    branch  
WHERE    assets>2500000 AND city='Edmonton'
```

- Simple predicates can be combined using AND, OR, NOT.

# Querying two Relations

**customer**(cname, street, city)

**deposit**(accno, cname, bname, balance)

- List the name and the city of every customer who has an account with balance over \$2,000.

```
SELECT customer.cname, city
FROM    customer, deposit
WHERE   balance > 2000
AND     customer.cname = deposit.cname
```

# Queries With Tuple Variables

**loan**(accno, cname, bname, amount)

**deposit**(accno, cname, bname, balance)

- Find customers who have both loans and deposits.

```
SELECT  loan.cname
FROM    loan, deposit
WHERE   loan.cname = deposit.cname
```

- Equivalently using tuple variables:

```
SELECT  l.cname
FROM    loan l, deposit d
WHERE   l.cname = d.cname
```

- Range variables are really needed if the same relation appears twice in the FROM clause.

# A Simple Evaluation Alg.

```
SELECT ...  
FROM   R1 r1, R2 r2, ...  
WHERE  C
```

- *Tuple variables*  $r_1, r_2, \dots$  respectively range over rows of  $R_1, R_2, \dots$
- Evaluation strategy:
  - FROM clause produces Cartesian product of listed tables
  - WHERE clause produces table containing only rows satisfying condition
  - SELECT clause retains listed columns

```
for every tuple  $r_1$  in  $R_1, r_2$  in  $R_2, \dots$   
  let  $r := r_1, r_2, \dots$   
  if  $C(r)$  then  
    output the desired columns
```



# From SQL to Relational Algebra

```
SELECT l.cname  
FROM   loan l, deposit d  
WHERE  l.cname = d.cname
```

Equivalent to:

$$\begin{aligned} temp &= loan \times deposit [a1, cn1, bn1, b1, a2, cn2, bn2, a2] \\ \pi_{cn1} \sigma_{cn1 = cn2} (temp) \end{aligned}$$

This is a simple evaluation algorithm for SELECT.

# Queries With Set/Bag Results

- Find all cities of customers.
  - **SELECT** city  
    **FROM** customer
- Result?
- To get rid of duplicates, we need
  - **SELECT** **DISTINCT** city  
    **FROM** customer

# Duplicates

- Duplicate rows not allowed in a relation
- However, duplicate elimination from query result is costly and not automatically done; it must be explicitly requested:

```
SELECT DISTINCT .....  
FROM .....
```

# Working with Strings

- Equality and comparison operators apply to strings (based on lexical ordering)
  - E.g. WHERE *cname* < 'P'
- Concatenate operator applies to strings
  - E.g. WHERE *bname* || '--' || *address* = ....
- Expressions can also be used in SELECT clause
  - E.g. SELECT *bname* || '--' || *address* AS NameAdd  
FROM branch

# Partial Matching

**customer**(cname, street, city)

- Find every customer whose address starts with “Computing Science”.

```
SELECT  *  
FROM    customer  
WHERE    address LIKE 'Computing Science%'
```

- Expression: col-name [NOT] LIKE pattern
  - Pattern may include wildcard characters ‘%’ matching any string and ‘\_’ (underscore) matching any single character.

# Naming the Results

**deposit**(accno, cname, bname, balance)

- For every deposit holder who has over \$1000, find the customer name and the balance over \$1000.

```
SELECT  cname, (balance - 1000) as bal
FROM    deposit
WHERE    balance > 1000
```

# Ordering the Results

**branch** (bname, address, city, assets)

- Find the names and assets of all branches with assets greater than \$2,500,000 and order the result in ascending order of asset values.

```
SELECT  bname, assets
FROM    branch
WHERE    assets > 2500000
ORDER BY assets
```

- Default is ascending order; a descending order can be specified by the DESC keyword.

# Queries Involving Set Operators

- set union : Q1 UNION Q2
  - the set of tuples in Q1 or Q2
- set difference : Q1 EXCEPT Q2
  - the set of tuples in Q1 but not in Q2
- set intersection : Q1 INTERSECT Q2
  - the set of tuples in both Q1 and Q2
- Q1 and Q2 must be union-compatibles
  - same number/types of attributes.



# Queries With Set Operators

- List deposit-holders who have no loans.

```
(SELECT  cname
FROM    deposit)
```

**EXCEPT**

```
(SELECT  cname
FROM    loan)
```

- List cities where there is either a customer or a branch.

```
(SELECT  city
FROM    customer)
```

**UNION**

```
(SELECT  city
FROM    branch)
```

# Queries With Set Operators

- Find all cities that have both customers and branches in.

```
(SELECT city  
FROM      customer)
```

**INTERSECT**

```
(SELECT city  
FROM      branch)
```

- Find every city that has a branch but no customer.

# Queries Over Multiple Relations

**branch** (bname, address, city, assets)

**customer**(cname, street, city)

**deposit**(accno, cname, bname, balance)

```
SELECT branch.bname, assets
FROM    branch, customer, deposit
WHERE   customer.city = 'Jasper'
AND     customer.cname = deposit.cname
AND     deposit.bname = branch.bname
```

- What does the query do?

Find the name and asset of every branch that has a deposit account holders who lives in Jasper.

# Queries With Nested Structures

- Queries within the WHERE clause of an **outer** query

```
SELECT  
FROM  
WHERE  OPERATOR (SELECT ... FROM ... WHERE)
```

- There can be multiple levels of nesting
- Operators: **IN**, **(NOT) EXISTS**, **< ALL**, ...
- **Avoid nesting as much as possible.**

# Nested Structures Using “IN”

**SELECT**

**FROM**

**WHERE** *expr* | (*expr list*) **IN**  
(**set of values**)

- E.g.

... **WHERE** *province* **IN** ( 'AB', 'BC' )

... **WHERE** *province* **NOT IN** ( 'AB', 'BC' )

# Example

**deposit**(accno, cname, bname, balance)

```
SELECT cname
FROM    deposit
WHERE    bname    IN
                (SELECT    bname
                 FROM    deposit
                 WHERE    cname = 'John Doe' )
```

- What does the query do?

Find every customer who has a deposit in some branch at which John Doe has a deposit.

# The Same Example Without Nesting

**deposit**(accno, cname, bname, balance)

```
SELECT    d1.cname
FROM      deposit d1, deposit d2
WHERE     d2.cname = 'John Doe'
AND      d1.bname = d2.bname
```

NOTE: avoid nesting as much as possible.

# Nested Structures Using “< ALL”, ...

```
SELECT  
FROM  
WHERE expr < ALL (set of values)
```

- Other forms: “<= ALL”, “= ALL”, “>= ALL”, “> ALL”
- Op ALL (set of values) evaluates to **true** iff the comparison evaluates to true for every value in the set.
- Op SOME (set of values) evaluates to **true** iff the the comparison evaluates to true for at least one value in the set.



# Nested Structures Using “>ALL”, ...

**branch** (bname, address, city, assets)

- Find branches that have assets greater than the assets of all branches in Calgary.

```
SELECT  bname
FROM    branch
WHERE   assets > ALL
        (SELECT assets
         FROM    branch
         WHERE   city = 'Calgary')
```

# Nested Structures Using “EXISTS”

**SELECT**

**FROM**

**WHERE (NOT) EXISTS (SELECT ... )**

- **EXISTS (SELECT ...)**  
evaluates to true iff the result of the subquery contains at least one row.
- The expression is evaluated for every iteration of the outer query.

# “EXISTS” Construct Example

**branch** (bname, address, city, assets)

**customer**(cname, street, city)


- Find the names of customers who live in a city with no bank branches.

```
SELECT  cname
FROM    customer
WHERE   NOT EXISTS (SELECT *
                    FROM    branch
                    WHERE    customer.city=branch.city)
```

- Find the names of customers who live in a city which has a bank branch.
  - Change NOT EXISTS to EXISTS.  
(can also write it using join)

# Division

- *Query type*: Find the subset of items in one set that are related to *all* items in another set
- *Example*: Find customers who have deposit accounts in *all* branches.



$$\pi_{cname, bname}(\text{deposit}) / \pi_{bname}(\text{branch})$$

# Division

- *Strategy for implementing division in SQL:*
  - Find set,  $A$ , of all branches in which a particular customer,  $c$ , has a deposit account.
  - Find set,  $B$ , of all branches.
  - Output  $c$  if  $A \subseteq B$ , or, equivalently, if  $B - A$  is empty

# Division – SQL Solution

**branch** (bname, address, city, assets)

**deposit**(accno, cname, bname, balance)

```
SELECT c.cname
FROM customer c
WHERE NOT EXISTS
    (SELECT b.bname                -- set B of all branches
      FROM branch b
      EXCEPT
      SELECT d.bname                -- set A of branches in which
                                   -- customer c has a deposit account
      FROM deposit d
      WHERE d.cname=c.cname        -- global variable
    )
```

# Division – 2<sup>nd</sup> SQL Solution

- Find customers who have deposit accounts in all branches.
  - Same as “find all customers such that there is no branch where they do not have deposits in.”

```
SELECT c.cname  
FROM   customer c  
WHERE NOT EXISTS
```

```
    (SELECT   bname  
     FROM     branch b  
     WHERE NOT EXISTS
```

```
        (SELECT *  
         FROM   deposit  
         WHERE  b.bname=deposit.bname  
         AND   c.cname = deposit.cname))
```

Deposits of c in  
branch b

branches where  
c has no deposits in

# Division – Another Example

- Find professors who have taught courses in *all* departments.
- *Strategy for implementing division in SQL:*
  - Find set,  $A$ , of all departments in which a particular professor,  $p$ , has taught a course
  - Find set,  $B$ , of all departments
  - Output  $p$  if  $A \subseteq B$ , or, equivalently, if  $B - A$  is empty



# Division – SQL Solution

Professor (*Id, Name, DeptId*)

Department (*DeptId, Name*)

Course (*DeptId, CrsCode, CrsName, Descr*)

Teaching (*ProfId, CrsCode, Semester*)

```
SELECT P.Id
FROM Professor P
WHERE NOT EXISTS
    (SELECT D.DeptId                -- set B of all dept Ids
     FROM Department D
     EXCEPT
     SELECT C.DeptId                -- set A of dept Ids of depts in
                                     -- which P has taught a course
     FROM Teaching T, Course C
     WHERE T.ProfId=P.Id          -- global variable
           AND T.CrsCode=C.CrsCode)
```

# SQL and SQLite

- SQLite: pretty modern but light language
  - Released in 2000 (Oracle was 1977, DB2 was 1983)
  - It is light (intended for small devices)
  - Tightly integrated with applications (software library rather than a stand-alone system/process to communicate)
  - Locks the whole file/database during writes (not a good choice for write-intensive/concurrent transactions)
  - Supports both in-memory and on-disk databases
- Differences with SQL-92
  - Does not support op ALL in nested queries
  - Does not support RIGHT OUTER JOIN and FULL OUTER JOIN
  - Views are read-only
  - ALTER TABLE command is very limited
  - Foreign keys constraints are not enforced by default
  - More left to be explored!

# Basic Data Types in SQLite

- **INTEGER** 1, 2, 3, 4, 6, 8 bytes integer  
(depending on the magnitude)
- **REAL** 8 bytes floating point number
- **TEXT** stored using database encoding (e.g. UTF-8)
- **BLOB** stored as is
- **NULL** null value
- **NUMERIC** can store all other types; stored value is converted to  
INTEGER or REAL if the conversion is lossless and reversible

# Other Types in SQLite

- |                       |         |
|-----------------------|---------|
| ■ INT, SMALLINT       | INTEGER |
| ■ CHAR(n), VARCHAR(n) | TEXT    |
| ■ DOUBLE, FLOAT       | REAL    |
| ■ DECIMAL             | NUMERIC |
| ■ DATE, DATETIME      | NUMERIC |

# SQL and Oracle

- Oracle
  - the first commercial relational database system.
  - based on the system R prototype.
- Will be using Oracle 11g.
- Many supporting products:
  - SQL Plus: interactive SQL interface
  - SQL Precompilers (C, C++)
  - SQL Loader: tool for mass loading of data
  - PL/SQL: Oracle's **P**rocedural **L**anguage extension
  - JDBC support
  - cx\_oracle: Python interface
  - Oracle Forms, Oracle Report Writer, etc.

# Basic Data Types in Oracle

- **NUMBER(m,n)** decimal number (precision, scale)  
(no parameters means floating point)
- **DATE** century,year,month,day, hh.mm.ss
- **CHAR(n)** fixed-length char string (n <= 2000)
- **VARCHAR2(n)** var-length char string (n <= 4000)
- **LONG** var-length char string (up to 2 gigabytes)
- **RAW(n)** fixed-length raw binary data (n <= 2000)
- **LONG RAW** var-length raw binary data (max 2 gigs)

# Other Data Types in Oracle

- VARCHAR(n) same as VARCHAR2(n)
- LONG VARCHAR(n) same as LONG
- DECIMAL(m,n) same as NUMBER(m,n)
- FLOAT same as NUMBER
- INTEGER same as NUMBER(38,0)
- SMALLINT same as NUMBER(38,0)
- REAL same as NUMBER
- DOUBLE PRECISION same as NUMBER