

Database Tuning

Measure of Performance

- **Response time** – average time to wait for a response to a particular query
- **Throughput** – volume of work completed in a fixed amount of time (often measured as transactions per second)

Normalization

- Can improve performance
 - Less redundancy
 - => more rows/page => less I/O
 - More tables
 - => smaller and more clustered indexes

Normalization (cont.)

- Can decrease performance
- E.g.:
 - Transcript(StudId, CrsCode, Semester, Grade)
 - Students(StuId, *name*)
 - The relations are in BCNF, *but* ...
 - a join required to list names of students with A in 291

```
SELECT S.Name
FROM Student S, Transcript T
WHERE S.Id = T.StudId AND T.CrsCode = '291'
      AND T.Grade = 'A'
```

- ... and join is expensive

Denormalize

- Add attribute Name to Transcript

```
SELECT T.Name  
FROM Transcript T  
WHERE T.CrsCode = 'CS305' AND T.Grade = 'A'
```

- Join is avoided, but added redundancy.

Performance Bottlenecks

- A table can be a performance bottleneck if
 - it is heavily used, causing locking contention
 - it's index is deep (has many rows), increasing I/O
 - rows are wide, increasing I/O
- Solution: table partitioning

Horizontal Partitioning

- Split rows
 - Geographically (e.g., by state), organizationally (e.g., by department), active/inactive (e.g., current students vs. grads)
- Useful when accesses are confined to disjoint subsets of rows
- Advantages:
 - Spreads users out and reduces contention (particularly if tables can be spread over different devices)
 - Indexes have fewer levels (but only marginally)
 - Rows in a typical result set are concentrated in fewer pages
- Disadvantages:
 - Added complexity
 - Difficult to handle queries over all tables

Vertical Partitioning

- Split columns and replicate the key
- Useful when table has many columns and
 - it is possible to distinguish between frequently and infrequently accessed columns
 - different queries use different subsets of columns
- **E.g.:** Split *employee* table into
 - Columns related to *compensation* (sin, tax, benefits, salary) and columns related to *job* (sin, department, projects, skills).
 - Decomposition is lossless join (since *sin* is included in both tables), although a join is required.

Horizontal Partitioning in Oracle

- Methods: range, hash, list, etc.
- e.g.

Create table transcripts (

 StudId char(6),

 CrsCode char(8),

 Semester char(3),

 Grade int,

 Primary Key (StudId, CrsCode, Semester))

Partition BY LIST (Semester) (

 Partition p10 values ('F10', 'W10'),

 Partition p11 values ('F11', 'W11'),

 Partition rest values (default));

Indexes

- Advantages: fast table access (for search, insert, delete, join, constraint checking)
- Created
 - automatically on primary key
 - Manually on other columns


Index Covering

- An index covers a query if
 - the query can be computed from the index alone
- All attributes of the query must be included in a covering index

Index Covering

- **E.g.** Professor(DeptId, Name, ...)
 - Dense B+-tree index on (DeptId, Name) covers the query because ...


```
SELECT P.Name  
FROM Professor P  
WHERE P.DeptId = 'CS'
```



Index Covering

- **E.g.** Student(Id, Name, Address, ...)
 - Dense B⁺-tree index on (Id, Name, Address)

```
SELECT S. Id, S.Name  
FROM Student S  
WHERE S.Address = '1 Lake St'
```



Choosing a Clustered Index

- Choose clustered index to support:
 - Point queries with large result sets
 - A clustered index on Transcript with search key (CrsCode, Semester) supports queries requesting the Id's of all students in a particular class.
 - Range queries
 - A clustered index on Transcript with search key (Grade) supports queries requesting the Id's of all students with grades between B+ and A
 - ORDER BY clauses
 - Sequence of attributes in ORDER BY is a prefix of search key
 - Index-nested (get all rows satisfying a particular search key value) and sort-merge (avoid sort) joins.

Choosing a Clustered Index

- Do not choose a clustered index
 - If it is not needed for one of the above
 - If the search key attribute is frequently updated
 - Since it will be necessary to move rows frequently
 - Hash – from bucket to bucket
 - B⁺ tree – from one leaf page to another

Choosing an Unclustered Index

- Choose unclustered index to support
 - Queries with small result sets
 - Index-nested loop joins (when the indexed column is in the inner relation)
 - Index covering
 - Point queries with small result sets

Miscellaneous Hints

- If an attribute is unique, declare it so (query optimizer can use it in query planning)
 - Only one row can match a search or join attribute
- Adjust fill factor
 - To 100% if table is read-only
 - To a smaller value if table is dynamic
- Keep search keys small to flatten index
- Add unclustered indexes only when necessary

Tuning SQL

- Hints:
 - Avoid sorts
 - Use of DISTINCT, UNION, EXCEPT, ORDER BY, GROUP BY cause sorts. Avoid their use if possible.

Tuning SQL

- Hints (con't):
 - Beware of views since they may cause unnecessary joins
 - Consider restructuring a query – different formulations will have different costs depending on state of tables and indexes available.

Influencing the Query Optimizer

- Statistics: Used by optimizers to estimate cost of a query plan (based on size of result sets)
 - Table: number of rows, number of distinct values of an attribute, max and min attribute values
 - Index: depth, number of leaf pages, number of distinct search key values
 - Histograms of attribute values
 - Example: use unclustered index if histogram shows that number of rows with attribute value specified in query is small, else use scan
 - Must be periodically updated if table dynamic

More Levels of Tuning

- Application level:
 - Query: schema redesign, use of indexes, client vs. server processing (stored procedures)
 - Transaction: isolation level, code design
- System level:
 - Cache issues: cache size, binding, I/O size
 - Distribution of data across devices
 - Log management
- Hardware level:
 - Number of cpu's
 - Disk configuration (many small disks vs. a few large ones)
 - Backup (mirrored disks vs logs)
- Distribution:
 - Replication
 - Distributing data and processing